

# inetd enhancements

---

Arjun Bemarker

## About the project

---

### Overview

`inetd` enhancements to allow it to babysit general daemons, not just networking ones: the ability to keep daemons alive (even when they crash), start them after various specified events/intervals and specify scheduling priority of them.

In addition, a utility will be added to start/stop these daemons.

### Current State

In its current form, you can tell `inetd` to startup network daemons by editing `/etc/inetd.conf` and specifying an ingress port. You can also list program arguments that you want `inetd` to supply. When there is an incoming connection, `inetd` will startup the daemon with the arguments and supply the incoming data to it through `stdin`.

### Deliverables

1. Update existing `inetd` code to support the new feature set
2. Feature documentation & User Guide
3. New utility to start/stop daemons

**I am proposing** to let `inetd` start other daemons based on events which aren't purely network events and keep them alive based on newly added parameters given to `inetd.conf`

### Specifics

specific parameters to add to `inetd.conf` :

#### `keep_alive`

- `bool`
- repeatedly restart Whenever the program crashes (non zero exit), `syslogd(8)` will be notified.

#### `successful_exit`

- `bool`
- if true, the daemon will be restarted until it fails.
- if false, the daemon will be restarted until it succeeds

#### `network_state`

- `boolean`
- if true, the daemon will only (re)start when network is available

#### `path_state`

- `string`
- if true, the program will be alive as long as the given path exists
- if false, the program will be alive as long the path doesn't exist

#### `throttle_interval`

- `int` (some default value will be decided)
- time in seconds to wait between program restarts

#### `nice`

- `int`
- run the program at an altered scheduling policy

And a command line utility which will allow the user to make individual configuration files and start and stop them on demand.

### Milestones and Schedule

Following the [Google Summer of Code schedule](#), coding officially begins on June 13th. By this time, I should be familiar with my mentor and up to speed with the codebase from the "Community bonding period" (May 20 - June 12).

Until the first evaluation phase, I have 7 weeks. I predict this project to be finished in 9 - 10 weeks.

A 350 hour project means it will come to ~32hr work week. I don't have any other commitments for this summer, so this should be doable.

Here is a brief overview. Of course, it is likely to take longer than I anticipate and plan to, so I have left two weeks as a buffer.

#### Week 1(starting June 12)

Make changes to `inetd.conf` syntax so it can accept processes which aren't necessarily invoked by activity on a network port.

#### Week 2,3,4

Create and add support for a command line utility so it can add/remove daemons. This should deal with unix credentials and IPC with the daemon.

Make inetd start processes based on events other than network ones.

#### Week 5

add `keep_alive`

#### Week 6

add `successful_exit`

#### Week 7

add `network_state` and `path_state`

#### Week 8

add `throttle_interval` and nice support

#### Week 9

#### Week 10

### Similar Software

MacOS has [launchd](#) Linux has [systemd](#)

### Licensing

All features are inspired by [launchd](#). Since that is under APSL, there shouldn't be any licensing issues.

## About the project and NetBSD

---

### Have you rebuilt the kernel and the userland, either in full or in parts?

Yes, I have rebuilt the inetd binary and got the tests running through atf.

### How will your project integrate into NetBSD

`inetd` is ran at boot time by `/etc/rc`. I will be modifying this program, so it should start the daemons it is responsible for. `/etc/inetd.conf` will be used to use the new features

### What interfaces in NetBSD will your project use?

in order to monitor and restart services when they crash, i will use various system calls including fork with `execve` and system libraries including `signal.h`, `wait.h` to setup signal handlers in order to handle the child daemons dying.

in order to check network availability, the `getifaddrs(3)` interface will be used. Here I can check whether the network is up with the `ifa_flags` (these flags are similar to the `ifconfig` utility)

Simple `access(2)` function will be used to wait until a file is created.

`setpriority(2)` will be used to set the "niceness" of the daemon.

there already exists a concept of time tracking in the `inetd` codebase through `ratelimit.c:rl_time()`. So, for `set_interval`, this interface can be used.

### To what degree are you familiar with those interfaces?

I am relatively familiar with them. I have been playing around with them over the past month in preparation for this project.

I've also looked through Apple's launchd source code and feel comfortable with using them.

## About you

---